


1. (0 pts.) (a). Run the Bellman-Ford algorithm on the graph given below, starting at the given source s . Whenever there is a choice of edges to update, always pick the one that is lexicographically/alphabetically first. Show the dist array after each round of updates.




Solution: The initial dist array is given in the first row of the matrix given below; after i -th round it is given in the $(i + 1)$ -th row of the matrix, $1 \leq i \leq 5$.

$$dist = \begin{bmatrix} s & b & c & d & e & f \\ 0 & \infty & \infty & \infty & \infty & \infty \\ 0 & 4 & 9 & 17 & 8 & 1 \\ 0 & 4 & 7 & 17 & 8 & 1 \\ 0 & 4 & 7 & 15 & 8 & 1 \\ 0 & 4 & 7 & 15 & 8 & 1 \\ 0 & 4 & 7 & 15 & 8 & 1 \end{bmatrix}$$

(b). Draw the shortest-path tree.

Solution:



2. (0 pts.) You are given a directed graph $G = (V, E)$ and a vertex $s \in V$. Each edge e is assigned with a length $l(e)$, possibly with negative value. We know that there is no negative cycle in this graph, and that the only negative edges are the ones that leave the vertex s . That is, $l(s, v) < 0$ for all $(s, v) \in E$, and $l(u, v) > 0$ for all $u \neq s$ and $(u, v) \in E$. If we run Dijkstra's algorithm starting at s , will it fail on this graph? Prove your conclusion.

Solution: No, it won't fail on this graph.

Proof: The correctness of Dijkstra's algorithm depends on the claim that the next closest vertex, i.e., v_{k+1}^* , must be within one-edge extension of R_k i.e. $v_{k+1}^* = \arg \min_{v \notin R_k, u \in R_k, (u, v) \in E} (\text{distance}(s, u) + l(u, v))$. The proof of this statement only uses that, the edges leaving any $v \notin R_k$ have positive edge length. In other words, the proof only requires that the one-edge extension is always preferred than the two-edge extension (first edge being (u, v) , second edge being (v, w) for some w) will always be longer since the second-edge has positive edge length. In our case, the second-edge always has positive length. This is because, although edges leaving s have negative edge length, s will always stay in R_k and consequently edges leaving s will always be part of one-edge extension.

3. (0 pts.) You are given a directed graph (V, E) with positive edge length $l(e)$ for any $e \in E$, and a positive weight $w(v)$ for any $v \in V$. Now we define the length of a path p from u to v as the sum of the lengths of all the edges in p plus the sum of the weights of all the vertices in p . You are also given a source $s \in V$ and it happens that $w(s) = 0$. Design an algorithm to find the length of the shortest path (in this new definition) from s to all vertices in V . Your algorithm should run in $O((|V| + |E|) \log |V|)$ time.

Solution: We create a new graph G' . For each vertex $v \in V$, we add two vertices v_i and v_o to G' , and add a new edge (v_i, v_o) with length being the weight of vertex v , i.e., $w(v)$. For each edge $e = (u, v)$ in G , we add an edge (u_o, v_i) to G' and its length keeps the same, i.e., $l(e)$. Notice that in G' nodes are not associated with weights and only edges are associated with lengths.

Clearly, any path $p = s \rightarrow w \rightarrow \dots \rightarrow u \rightarrow v$ in G corresponds to its counterpart $p' = s_i \rightarrow s_o \rightarrow w_i \rightarrow w_o \rightarrow \dots \rightarrow u_i \rightarrow u_o \rightarrow v_i \rightarrow v_o$. And the length of p , under the new definition (i.e., sum of vertex weights and edge lengths), is exactly the length of p' in G' , under the regular definition (i.e., sum of edge lengths). Therefore, we can run Dijkstra's algorithm on G' , with s_i being the source vertex. The distance from s_i to v_o in G' will give exactly the length of the shortest path from s to v in G (under new definition).

Graph G' contains $2|V|$ vertices and $(|V| + |E|)$ edges. The Dijkstra's algorithm, which dominates the running time, takes $O(2|V| + |V| + |E|) \cdot \log(2|V|) = O(|V| + |E|) \cdot \log(|V|)$ time.

4. (0 pts.) Shortest path algorithms can be applied in currency trading. Let c_1, c_2, \dots, c_n be various currencies. For any two currencies c_i and c_j , there is an exchange rate $r_{i,j}$; this means that you can purchase $r_{i,j}$ units of currency c_j in exchange for one unit of c_i . These exchange rates satisfy the condition that $r_{i,j} \cdot r_{j,i} < 1$, so that if you start with a unit of currency c_i , change it into currency c_j and then convert back to currency c_i , you end up with less than one unit of currency c_i (the difference is the cost of the transaction).

(a). Give an efficient algorithm for the following problem: given a set of exchange rates $r_{i,j}$, and two currencies s and t , find the most advantageous sequence of currency exchanges for converting currency s into currency t . Toward this goal, you should represent the currencies and rates by a graph whose edge lengths are real numbers.

Solution. We can transform the currency exchange problem into the shortest path problem. We build a directed graph $G = (V, E)$, where $V = \{c_1, c_2, \dots, c_n\}$ represents all currencies, and E contains all pairs (c_i, c_j) , $1 \leq i \neq j \leq n$. We assign length for edge $(c_i, c_j) \in E$ as $-\log(r_{i,j})$. We

now show that computing the shortest path from s to t in G actually gives the optimal sequence of currency exchanges for converting s into t . In fact, there is one-to-one correspondence between a path from s to t in G and a sequence of currency exchange for converting s into t . Moreover, the length of such a path p equals to $\sum_{(c_i, c_j) \in p} -\log r_{i,j} = -\log \prod_{(c_i, c_j) \in p} r_{i,j}$. Hence, the shortest path in G gives the path maximizes $\prod_{(c_i, c_j) \in p} r_{i,j}$, which is exactly the maximized amount of currency t that can be converted from a unit of currency s .

Based on the above analysis, the algorithm will be to compute the shortest path from s to t in G (using any of the single-source shortest path algorithm we introduced). The vertices along this optimal path gives the sequence of currencies following which we can get the maximized amount of currency t .

(b). The exchange rates are updated frequently. Occasionally the exchange rates satisfy the following property: there is a sequence of currencies c_1, c_2, \dots, c_k such that $r_{1,2} \cdot r_{2,3} \cdot r_{3,4} \cdots r_{k-1,k} \cdot r_{k,1} > 1$. This means that by starting with a unit of currency c_1 and then successively converting it to currencies c_2, c_3, \dots, c_k , and finally back to c_1 , you would end up with more than one unit of currency c_1 . Give an efficient algorithm for detecting the presence of such an anomaly. Use the graph representation you found above.

Solution. We need to detect whether there is $r_{1,2} \cdot r_{2,3} \cdot r_{3,4} \cdots r_{k-1,k} \cdot r_{k,1} > 1$, which is equivalent to detect $-(\log(r_{1,2}) + \log(r_{2,3}) + \log(r_{3,4}) + \cdots + \log(r_{k-1,k}) + \log(r_{k,1})) < 0$. Since we assign length for edge $e = (c_i, c_j)$ as $-\log(r_{i,j})$, this exactly implies a negative cycle in G . In other words, such an anomaly exists if and only if G contains negative cycles. Therefore, we can use Bellman-Ford algorithm on G to identify whether G contains negative cycles, which also answers whether there exists an anomaly.