

The Max-Flow Min-Cut Theorem

We now prove that the Ford-Fulkson algorithm is correct, i.e., when the algorithm terminates, the returned flow is a maximum-flow. Let f^* be the flow when the FF terminates. Note that the residual graph G_{f^*} w.r.t. f^* does not have any $s-t$ path (or equivalently, s cannot reach t in G_{f^*}), because otherwise the FF algorithm will not terminate. Now we define a cut: let S^* be the set of vertices in G_{f^*} that can be reached from s , formally $S^* := \{v \in V \mid s \text{ can reach } v \text{ in } G_{f^*}\}$, and define $T^* = V \setminus S^*$. Clearly, $s \in S^*$. Since s cannot reach t in G_{f^*} , we know that $t \in T^*$. Hence (S^*, T^*) is an $s-t$ cut. For example, in Figure 4, $S^* = \{s, b\}$ and $T^* = \{a, t\}$.




Figure 1: Illustration the proof of the max-flow min-cut theorem. Left: the residual graph G_{f^*} wrt the flow f^* when FF terminates (see Figure 4 of Lecture 19). The cut (S^*, T^*) is constructed by collecting vertices reachable from s in G_{f^*} as S^* , and $T^* = V \setminus S^*$. Right: the same $s-t$ cut (S^*, T^*) drawn on the (original) graph G , where one can see that for every edge $e \in E(S^*, T^*)$, $f^*(e) = c(e)$, and for every edge $e \in E(T^*, S^*)$, $f^*(e) = 0$.

We now show the central result: $|f^*| = c(S^*, T^*)$. Recall the Fact 2 in Lecture 18: the value of any flow, here we focus on f^* , can be calculated with any $s-t$ cut, and here we will use the particular cut (S^*, T^*) . Specifically, we can write $|f^*| = \sum_{e \in E(S^*, T^*)} f^*(e) - \sum_{e \in E(T^*, S^*)} f^*(e)$.

In fact, we have $f^*(e) = c(e)$ for every $e \in E(S^*, T^*)$. See Figure 1. Why? Suppose conversely that $f^*(e) < c(e)$. Assume that $e = (u, v)$; the fact that $e = (u, v) \in E(S^*, T^*)$ means that $u \in S^*$ and $v \in T^*$. According to the way the residual graph is constructed, in G_{f^*} , there will be a forward edge (u, v) . Consequently, s can reach v in G_{f^*} , as s can reach u (because $u \in S^*$) and u can reach v following this forward edge. This contradicts to the fact that $v \in T^*$.

Also, we can show that $f^*(e) = 0$ for every $e \in E(T^*, S^*)$. Why? Suppose conversely that $f^*(e) > 0$. Assume that $e = (u, v)$; the fact that $e = (u, v) \in E(T^*, S^*)$ means that $u \in T^*$ and $v \in S^*$. Then in the residual graph G_{f^*} , there will be a backward edge (v, u) . Consequently, s can reach u in G_{f^*} , as s can reach v (because $v \in S^*$) and v can reach u following this backward edge. This contradicts to the fact that $u \in T^*$.

Combining all above, we have $|f^*| = \sum_{e \in E(S^*, T^*)} f^*(e) - \sum_{e \in E(T^*, S^*)} f^*(e) = \sum_{e \in E(S^*, T^*)} c(e) - 0 = c(S^*, T^*)$.

This proves that f^* is a maximum-flow and (S^*, T^*) is a minimum $s-t$ cut, i.e., the FF algorithm is optimal, and a minimum-cut can be constructed from the residual graph. It also completes the connection between the maximum-flow problem and the minimum-cut problem: for any network, the value of the maximum-flow always equal to the capacity of the minimum-cut. This is called the max-flow min-cut theorem.

We analyze the time complexity of the Ford-Fulkson Algorithm. If we assume that all capacities are integers, we have $x(p) \geq 1$ and therefore the value of the new flow will be increased by at least 1 in each iteration. Hence, the number of iterations is upper-bounded by $|f^*|$. Each iteration includes building G_f , finding an $s-t$ path p in G_f (with BFS/DFS), and augmenting p , which takes $O(|V| + |E|)$ time. Therefore, the FF algorithm runs in $O(|f^*| \cdot (|V| + |E|))$ time.

Maximum-Matching and Minimum-Vertex Cover

Let $G = (V, E)$ be an undirected graph. A *matching* $M \subset E$ is a subset of edges of G that does not have common vertices. A matching is also called an *independent edge set*. See Figure 2 for an example. Given an undirected graph G , the so-called *maximum-matching problem* seeks a matching M of G such that the number of edges in the matching, i.e., $|M|$, is maximized.




Figure 2: A matching M and a vertex cover V_1 of an undirected graph G . Note that M is also a *maximum matching* of G and V_1 is also a minimum vertex cover of G .

A *vertex cover* of an undirected graph $G = (V, E)$ is defined as a subset of vertices $V_1 \subset V$ that *covers* all edges, i.e., for every edge $(u, v) \in E$ either $u \in V_1$ or $v \in V_1$ or both. Given an undirected graph G , the so-called *minimum-vertex cover problem* seeks a vertex cover V_1 of G such that the number of vertices in it, i.e., $|V_1|$, is minimized.

There exists a connection between matchings and vertex covers of an undirected graph, stated below.

Claim 1. Let $G = (V, E)$ be an undirected graph. Let M be any matching of G and let V_1 be any vertex cover of G . Then we have $|M| \leq |V_1|$.

Proof. Since V_1 is a vertex cover, it covers all edges of G , in particular, all edges of M . Since M is a matching, all edges in M does not share vertices, so it requires $|M|$ vertices to cover edges in M . Combined, we have $|M| \leq |V_1|$. \square

So, the size of any vertex cover is an upper bound of the size of any matching. We again can visualize such relationship with Figure 3. Clearly, if we can find a matching M and a vertex cover V_1 (of the same undirected graph G) satisfying that $|M| = |V_1|$, then M must be a maximum matching and V_1 must be a minimum vertex cover.




Figure 3: Each blue point represents the size of a matching and each red point represents the size of a vertex cover, of the same undirected graph. Do they always meet?

Does for *every* undirected graph there always exists a matching M and a vertex cover V_1 that satisfies $|M| = |V_1|$? The answer is no, and Figure 2 gives such an example. A more informative example is a triangle: a graph $G = (V, E)$ with $V = \{a, b, c\}$ and $E = \{(a,b), (b,c), (c,a)\}$. In a triangle, the maximum matching includes one edge, but to cover all 3 edges we need 2 vertices. In fact, an undirected graph consists of

an *odd-length cycle* with $2m + 1$ edges admits a maximum matching with m edges and a minimum vertex cover with $m + 1$ vertices. It seems that it's the odd-length cycles that results in a gap. In fact, this is true. Following we will show that, if an undirected graph G does not contain odd-length cycle, then the size of its maximum matching must be equal to the size of its minimum vertex cover. The proof is again constructive: we will design an algorithm to find both (the algorithm is an application of max-flow), and show that they have equal size.

We give an equivalent (and nice) characterization of an undirected graph without odd-length cycles.

An *bipartite* graph, usually written as $G = (X \cup Y, E)$, is an undirected graph whose vertices can be separated into two disjoint subsets X and Y such that all its edges must span X and Y (i.e., every edge $e = (u, v) \in E$ must satisfy $u \in X$ and $v \in Y$). See Figure 4.




Figure 4: A bipartite graph $G = (X \cup Y, E)$, a maximum matching M , and a minimum vertex cover V_1 of G .

Fact 1. An undirected graph G does not contain odd-length cycle if and only if G is a *bipartite* graph.

Proof. We first prove that if G is bipartite then any cycle in it must be even-length. (See cycle (b, h, c, f) in Figure 4.) This is because any cycle need to visit X and Y alternately so its length must be even.

For the other side, assume that every cycle of undirected graph G is even and we now prove that G is




Figure 5: The shortest path tree T where vertices are assigned to X or Y based on the parity of the distance from s . Left: an edge $(u, v) \in E$ is part of the tree. Right: an edge $(u, v) \in E$ is not part of the tree.

bipartite. We assume that G consists of one connected component; otherwise all connected components of G can be processed separately in the same way below. Let s be any vertex in G . Let $distance(s, v)$ as the length of the shortest path from s to v . Let T be the shortest path tree wrt s , i.e., the number of edges on the path from s to v is exactly $distance(s, v)$. See Figure 5. We now put $v \in V$ into X if $distance(s, v)$ is even, and put $v \in C$ into Y if $distance(s, v)$ is odd. Clearly, X and Y is a partition of all vertices. We now show that, any edge $(u, v) \in E$ must span X and Y . There are two cases. The first one is that (u, v) is also an edge in T . In this case we must have that $distance(s, v) = distance(s, u) + 1$ so u and v must be assigned differently. (See Figure 5, Left.) The second case is that $(u, v) \notin T$. Suppose conversely that $distance(s, u)$ and $distance(s, v)$ have the same parity. Consider the $s-u$ path and the $s-v$ path in T . Let b be the branching point of these two paths (also called the least common ancestor of u and v). Then the $b-u$ path and the $b-v$ path are edge-disjoint and their lengths must also have the same parity. Now these two paths combined with edge (u, v) must form a cycle of odd-length in G , a contradiction. (See Figure 5, Right.) \square